

Analysis of Graduate School Acceptance Chance

Final Report

Stat 4511

Group A

Group Members: Dagan Larson, Cole Rehbein, Shane Pratt

Due: April 29, 2025

Contents

1	Introduction	2
2	Regression Analysis	2
2.1	Model Fitting	2
2.2	Cross Validation	2
2.3	Assumptions and Remedial Measures	3
2.4	Final Model	3
2.5	Predictions from Final Model	4
2.6	Figures	5
3	Discussion	9
4	Conclusion	9
5	Additional Work	9
6	Appendix	10
6.1	Importing Data	10
6.2	Model Selection	10
6.2.1	Cross Validation	11
6.3	Model Assumptions	12
6.3.1	Polynomial Model Fitting	12
6.3.2	Identifying Outliers	15
6.4	Weighted Model	16
6.5	Bootstrapping	16
6.5.1	Confidence Intervals	16
6.5.2	Prediction Intervals	20

1 Introduction

2 Regression Analysis

2.1 Model Fitting

As there are only 7 variables in the dataset, model fitting was done exhaustively. Comparing these models via AIC , the best two models are as follows:

$$\text{Model A: } \hat{Y}_i \sim b_0 + b_1X_{i1} + b_2X_{i2} + b_3X_{i3} + b_4X_{i4} + b_5X_{i5}$$

$$\text{Model B: } \hat{Y}_i \sim b_0 + b_1X_{i1} + b_2X_{i2} + b_3X_{i3} + b_4X_{i4} + b_5X_{i5} + b_6X_{i6}$$

Where:

- \hat{Y}_i is estimated probability of acceptance
- X_{i1} is GRE score
- X_{i2} is TOEFL score
- X_{i3} is letter of recommendation strength
- X_{i4} is cumulative GPA
- X_{i5} is whether or not student did undergraduate research
- X_{i6} is university rating

They have AIC s of -1059.225 and -1058.386, respectively. These two models were also compared with BIC , adjusted R^2 , $PRESS$, and C_p . The results are summarized below:

Statistic	Model A	Model B
AIC	-1059.225	-1058.386
BIC	-1031.284	-1026.454
R_a^2	0.8002	0.8003
$PRESS$	0.7966	0.7963
C_p	5.494	6.353

As shown, Model A is preferred by AIC and BIC , whereas R_a^2 , $PRESS$, and C_p prefer Model B. So, cross validation was conducted to see which of the two models should be used.

2.2 Cross Validation

As there are only 400 observations in the dataset, leave-one-out CV was employed. We find that:

$$MSPE_A = .00413$$

$$MSPE_B = .00413$$

So, the two models have almost identical predictive ability. Additionally, the extra variable in Model B as compared to Model A, namely university rating, was found to not be statistically relevant ($p = 0.29$). Given these facts, we selected Model A for its simplicity.

2.3 Assumptions and Remedial Measures

To begin, the five predictor variables in Model A were examined for their linear relationship with Y . As shown in [Figure 1](#), there is linear relationship between Y and all of the predictors. Also present is significant multicollinearity between the predictors, however, nothing was done to remedy this as this model only seeks to make predictions in the scope of the data.

On the other hand, assumptions of error normalcy and error homoscedasticity are violated in this model, see [Figures 2](#) and [3](#). Transformations on Y were attempted to fix this, however, this alone was not sufficient. Examination of residual vs. predictor plots showed that polynomial terms for each of the continuous predictors may be necessary.

A hierarchical approach was taken to fitting these polynomial models, wherein a polynomial model featuring quadratic and cubic terms for GRE, TOEFL, and CGPA were introduced. These terms were iteratively deleted by taking the one with the highest p-value > 0.05 and removing it. If a lower-order term was removed this way, the higher power terms would also be removed. Centered variables were used for these power terms. This was repeated until all terms are relevant. Applying this approach also did not yield with acceptable error normalcy or distribution, so outliers were investigated for their impact on the model.

It was found that there are five data major outliers with respect to Y . When these five outliers were omitted from the dataset and the procedure outlined above was repeated on this new dataset. When this was done, a model with normal but heteroskedastic errors was constructed. A weighted least squares model was constructed to fix this heteroskedasticity, forming our final model which meets all of the assumptions of regression.

2.4 Final Model

The final model is:

$$(\hat{Y}_i)^5 \sim b_0 + b_{w1}X_{i1} + b_{w2}X_{i2} + b_{w3}X_{i3} + b_{w4}X_{i4} + b_{w5}X_{i5} + b_{w22}x_{i2}^2 + b_{w44}x_{i4}^2 + b_{w444}x_{i4}^3$$

Where:

- \hat{Y}_i is estimated probability of acceptance
- X_{i1} is GRE score
- X_{i2} is TOEFL score
- x_{i2} is centered TOEFL score ($x_{i2} = X_{i2} - \bar{X}_2$)

- X_{i3} is letter of recommendation strength
- X_{i4} is cumulative GPA
- x_{i4} is centered CGPA
- X_{i5} is whether or not student did undergraduate research
- b_{wi} 's are weighted regression coefficients

This model was constructed from the dataset omitting major Y outliers, namely points 66, 67, 69, 116, and 360. All of the coefficients are significant to a $\alpha = 0.05$ level.

The values of these coefficients are listed below alongside bootstrapped confidence intervals with a family wide confidence level of 90% (individual confidence level of $\alpha = 0.1/9$). See [Figure 4](#) for plots showing the estimated distributions of these coefficients.

Coefficient	Estimated Value	90% Family Confidence Limits
b_{w0}	-2.5315	(-2.9874 , -2.0727)
b_{w1}	0.0013	(-0.0003 , 0.0026)
b_{w2}	0.0053	(0.0025 , 0.008)
b_{w3}	0.0276	(0.0162 , 0.0401)
b_{w4}	0.1928	(0.1516 , 0.243)
b_{w5}	0.0337	(0.0102 , 0.0588)
b_{w22}	0.0003	(0 , 0.0005)
b_{w44}	0.1423	(0.119 , 0.1705)
b_{w444}	0.0357	(0.0048 , 0.0591)

2.5 Predictions from Final Model

Finally, we used the final model to generate two simultaneous predictions at a family-wide 95% level. We wanted to see how much an average (in the scope of the data) student's acceptance chance changes if they did or did not do research in their undergrad. We let:

- $\text{GRE}_{new} = 320$
- $\text{TOEFL}_{new} = 110$
- $\text{LOR}_{new} = 4.0$
- $\text{CGPA}_{new} = 8.6$
- $\text{Research} = 0$ for no research and 1 for research

These two new points were confirmed to be well within the range of our data and not extrapolation. We find that:

	Predicted Value (\hat{Y}_{new}^5)	95% family PIs (\hat{Y}_{new}^5)
Did Research	0.2651	(0.2471 , 0.2820)
No Research	0.2313	(0.2129 , 0.2487)

	Reverse-Transformed Prediction (\hat{Y}_{new})	Reverse-Transformed PIs (\hat{Y}_{new})
Did Research	0.7668	(0.7561 , 0.7763)
No Research	0.7462	(0.7339 , 0.7571)

See Figure 5 for distribution of predicted values. Note that the endpoints for the PIs were calculated before the reverse transformation was applied.

As shown in the prediction intervals, whether or not a student with average academics did research seems to only slightly increases their chance of acceptance into graduate school, if all else is held constant.

2.6 Figures

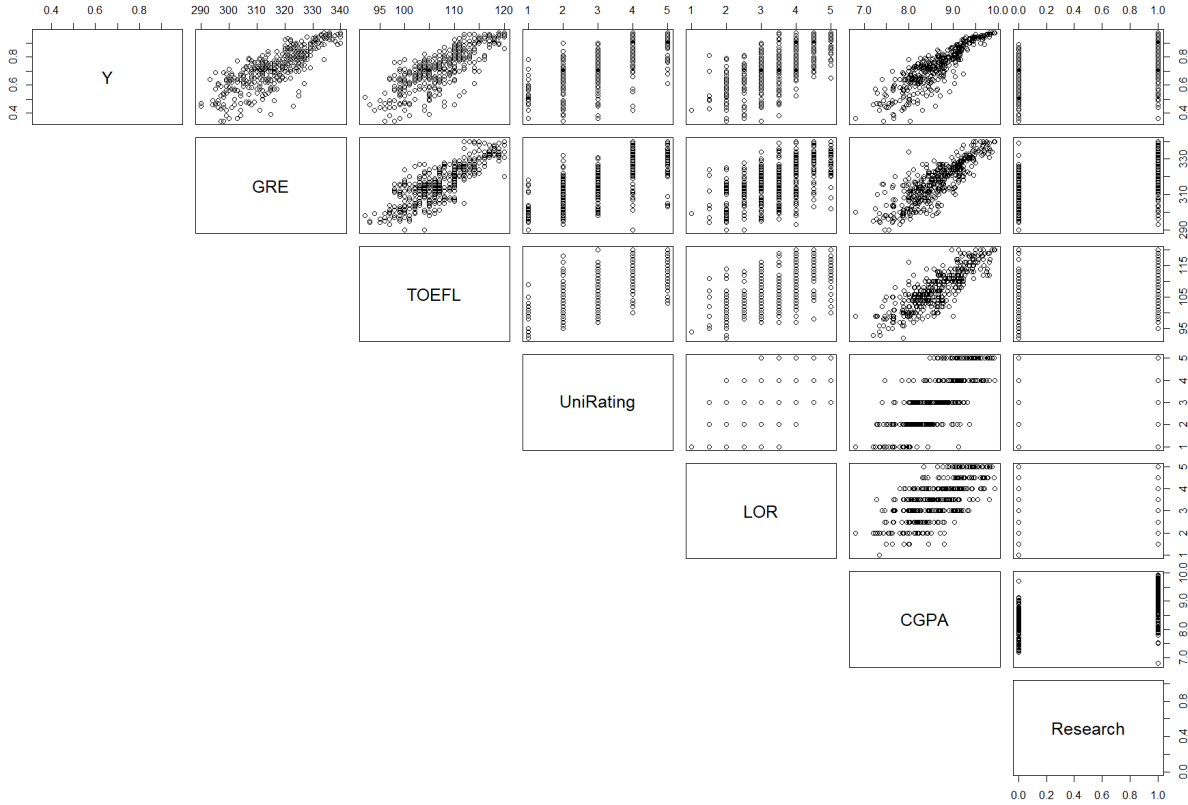


Figure 1: Plots of various predictor variables versus response and against each other. There is linear relationship with Y and all of the X 's. There is also significant multicollinearity present.

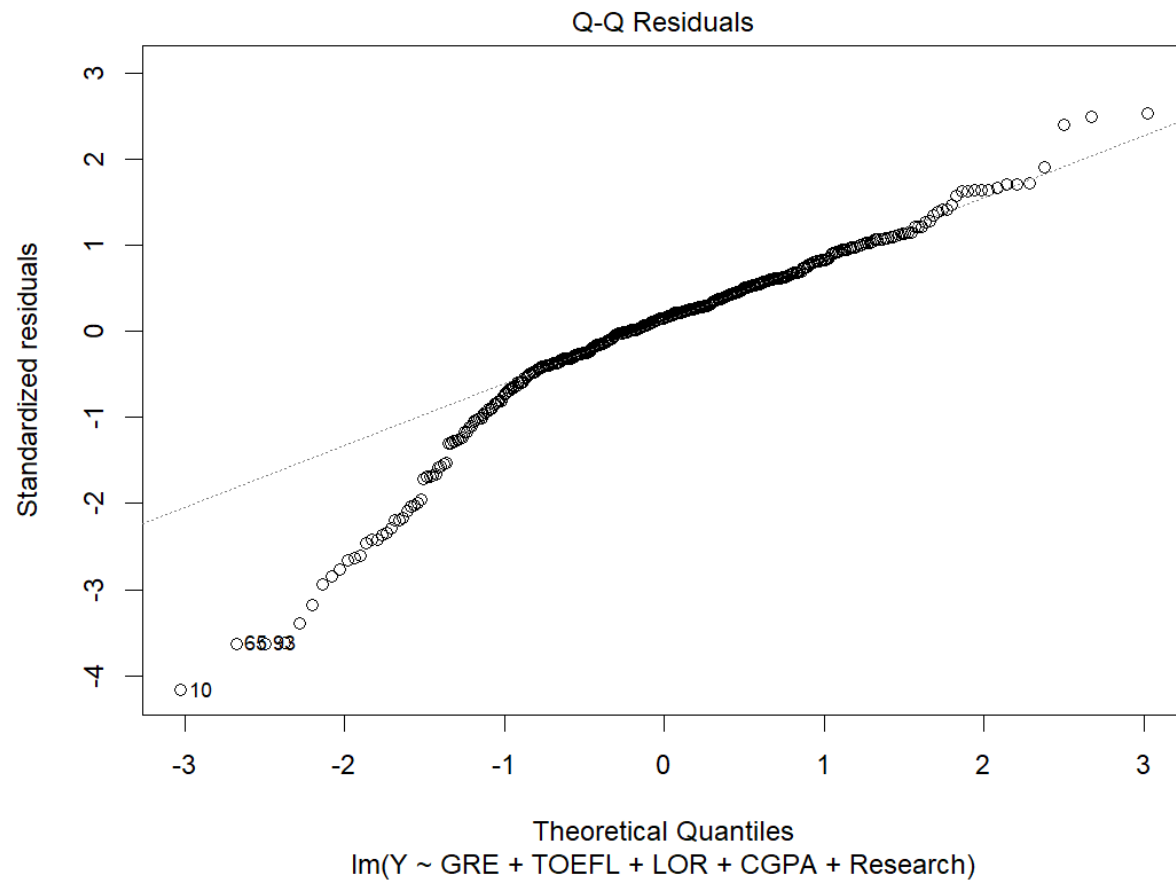


Figure 2: Normal probability plot of basic 5-variable model. Assumption of normality has been violated.




Figure 3: Error variance of basic 5-variable model, heteroskedasticity present.




Figure 4: Plots of bootstrapped weighted coefficients over 1000 trials

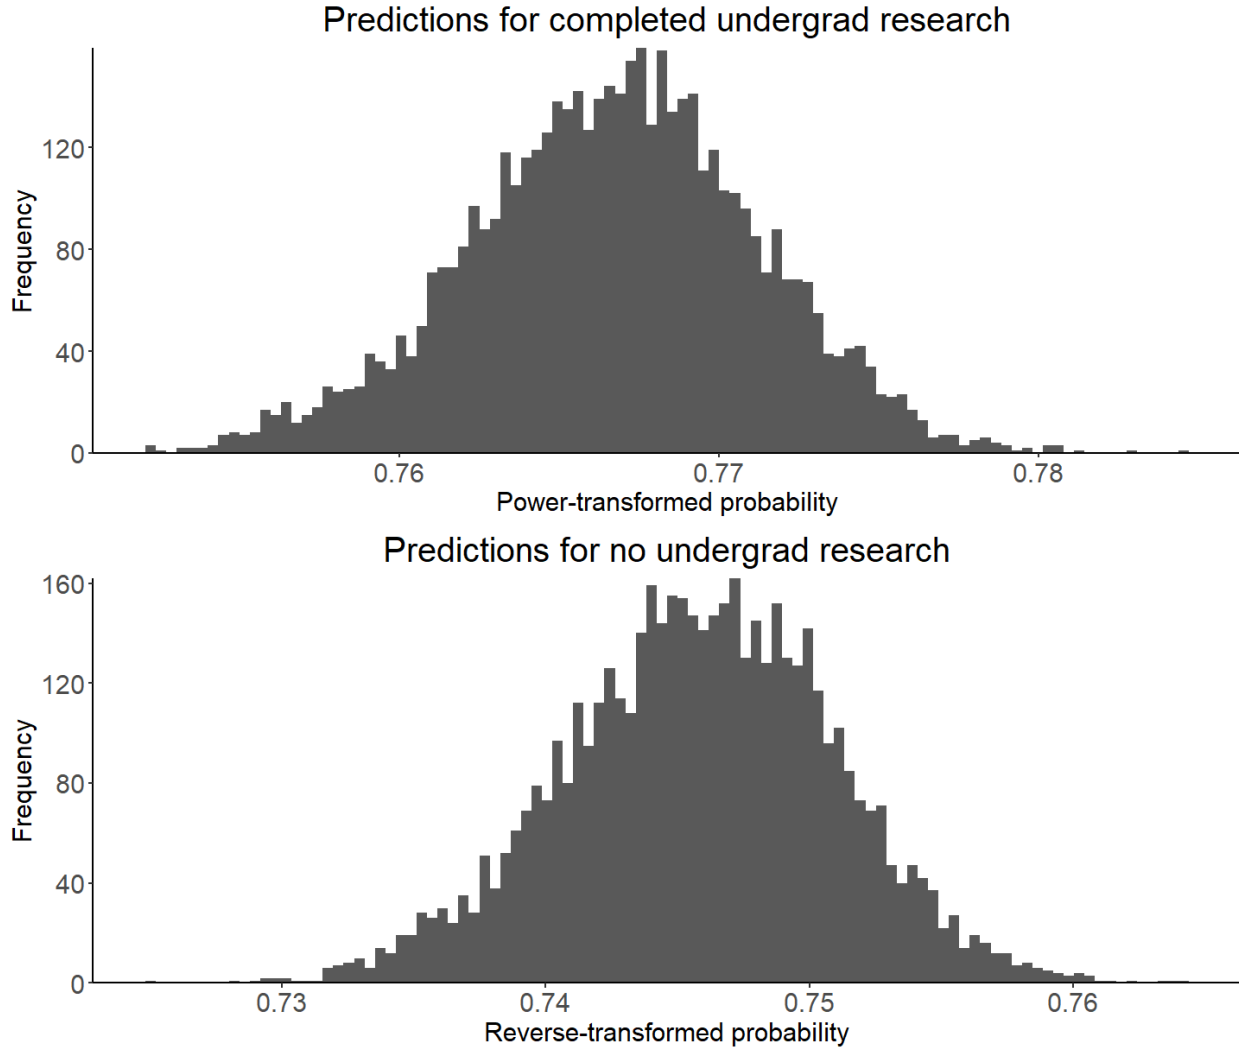


Figure 5: Plots of bootstrapped (5000 iterations) predicted probability of acceptance for an average student with or without having done undergraduate research. Values have already been transformed back into terms of the original data

3 Discussion

4 Conclusion

5 Additional Work

6 Appendix

6.1 Importing Data

```
1 Admission <- read.table("./Admission_Predict.csv", header = TRUE, sep = ",")
2
3 Admission = Admission[,which(names(Admission) != "Serial.No.")] # remove
  serial.no.
4 library(dplyr) # for rename
5 Admission = rename(Admission,
6                     c(GRE= "GRE.Score", TOEFL = "TOEFL.Score", UniRating = "
  University.Rating", Y = "Chance.of.Admit"))
7 Admission |> head()
```

6.2 Model Selection

```
1 library(ExhaustiveSearch)
2
3 # Exhaustive search via AIC
4 es_AIC = ExhaustiveSearch(formula = Y ~ ., data = Admission, family = '
  gaussian', performanceMeasure = "AIC")
5 print(es_AIC)
6 # Top models are:
7 # GRE + TOEFL + LOR + CGPA + Research
8 # AIC: -1059.225
9 # GRE + TOEFL + UniRating + LOR + CGPA + Research
10 # AIC: -1058.386
11
12 model_5var = lm(data = Admission, formula = Y ~ GRE + TOEFL + LOR + CGPA +
  Research)
13 model_6var = lm(data = Admission, formula = Y ~ GRE + TOEFL + UniRating + LOR
  + CGPA + Research)
14 summary(model_5var)
15 summary(model_6var)
16 # Adj R^2 of 5 var : 0.8002
17 # Adj R^2 of 6 var : 0.8003
18
19 library(lme4) # for BIC
20
21 BIC(model_5var) # -1031.284
22 BIC(model_6var) # -1026.454
23
24 library(qpcR) # for PRESS
25
26 PRESS(model_5var, verbose = FALSE) # 0.796567
27 PRESS(model_6var, verbose = FALSE) # 0.7962553
28
29 # PRESS for 6 var model is slightly higher than PRESS of 5 var model
30
31 library(olsrr) # for mallow's CP
```

```

32
33 full_model = lm(data = Admission, formula = Y ~ GRE + TOEFL + LOR + CGPA +
    Research + UniRating + SOP)
34
35 ols_mallows_cp(model_5var, full_model) # 5.494153
36 ols_mallows_cp(model_6var, full_model) # 6.353168
37
38 # 6 variable model slightly less biased
39
40 # So:
41 # AIC and BIC favor 5 variable model
42 # Adj R2, PRESS, and CP favor 6 variable model
43 # Cross validation necessary

```

6.2.1 Cross Validation

```

1 # CV will be conducted using leave-one-out method
2
3 TotalPE_5 = 0
4 TotalPE_6 = 0
5
6 N = 400
7
8 for(i in 1:N){
9   mod_5 = lm(data = Admission[-c(i),], formula = Y ~ GRE + TOEFL + LOR + CGPA +
    Research)
10  mod_6 = lm(data = Admission[-c(i),], formula = Y ~ GRE + TOEFL + LOR + CGPA +
    Research + UniRating)
11
12  PE_5 = Admission[i,"Y"] - predict(mod_5, Admission[i,])
13  PE_6 = Admission[i,"Y"] - predict(mod_6, Admission[i,])
14
15  TotalPE_5 = TotalPE_5 + PE_5^2
16  TotalPE_6 = TotalPE_6 + PE_6^2
17 }
18
19 MSPE_5 = TotalPE_5/N
20 MSPE_6 = TotalPE_6/N
21
22 round(MSPE_5, 5) # 0.00413
23 round(MSPE_6, 5) # 0.00413
24
25 # So, the 5 and 6 variable models have almost the exact same predictive power
26 # Given this and given that the extra variable (UniRating) in 6 variable model
27 # is not statistically significant (P-value = 0.29), 5 variable model should
    be selected
28 # for simplicity

```

6.3 Model Assumptions

```
1 library(car) # for ncvTest (Breusch-pagan)
2
3 base_model = lm(data = Admission, formula = Y ~ GRE + TOEFL + LOR + CGPA +
4   Research)
5 plot(base_model)
6 # residuals look nonnormal and heteroskedastic, no large residuals via Cook's
7   Distance however
8
9 shapiro.test(base_model$residuals) # not normal
10 ncvTest(base_model) # heteroskedasticity present
11
12 # Plots on predictors vs Y to detect violations of linearity between Y and
13   each predictor
14 pairs(Admission[c("Y", "GRE", "TOEFL", "UniRating", "LOR", "CGPA", "Research")],
15   lower.panel = NULL)
16 # Y appears to be linear wrt. each predictor, no violation there
17 # multicollinearity present, but the desire is to use model to make
18   predictions in scope
19 # of data, nothing will be done about it
20
21 # So, linearity and outlier assumptions seem fine, need to fix error
22   distrubtion
23 # Done using Boxcox transformation
24
25 library(EnvStats) # for Boxcox
26
27 bc = boxcox(base_model, lambda = seq(-6,10, 0.1))
28 plot(bc) # optimal value appears to be around 5 or 6
29 bc$lambda[which.max(bc$objective)] # 5.6
30 # Use Y^6 transform model
31
32 model_Y6 = lm(data = Admission, formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA +
33   Research)
34 plot(model_Y6)
35 shapiro.test(model_Y6$residuals) # not normal
36 ncvTest(model_Y6) # heteroskedasticity present
37
38 # Normality did get better, but heteroskedasticity is still bad, introduce
39   polynomial
40 # terms to model to try to fix this
```

6.3.1 Polynomial Model Fitting

```
1 # A hierarchical approach is taken, wherein a polynomial model featuring
2   quadratic and
3   # cubic terms for GRE, TOEFL, and CGPA are introduced. These terms are
4     iteratively deleted
```

```

3 # by taking the one with the highest p-value > 0.05 and removing it. If a
  lower-order term
4 # is removed this way, the higher power terms will also be removed. This will
  be repeated until
5 # all terms are relevant.
6 # Additionally, variables are centered for higher-order terms
7
8 poly_mod1 = lm(data = Admission,
9               formula = Y ~ GRE + TOEFL + LOR + CGPA + Research +
10                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((CGPA
11                  - mean(CGPA))^2) +
12                  I((GRE - mean(GRE))^3) + I((TOEFL - mean(TOEFL))^3) + I((CGPA
13                  - mean(CGPA))^3))
14 summary(poly_mod1)
15 # Least relevant term is (centered) TOEFL ^ 2, remove that and it's cubic term
16
17 poly_mod2 = lm(data = Admission,
18               formula = Y ~ GRE + TOEFL + LOR + CGPA + Research +
19                  I((GRE - mean(GRE))^2) + I((CGPA - mean(CGPA))^2) +
20                  I((GRE - mean(GRE))^3) + I((CGPA - mean(CGPA))^3))
21 summary(poly_mod2)
22 # Least relevant term is (centered) CGPA ^ 2, remove that and it's cubic term
23
24 poly_mod3 = lm(data = Admission,
25               formula = Y ~ GRE + TOEFL + LOR + CGPA + Research +
26                  I((GRE - mean(GRE))^2) + I((GRE - mean(GRE))^3))
27 summary(poly_mod3)
28 # Cubic term not relevant, remove that
29
30 poly_mod4 = lm(data = Admission,
31               formula = Y ~ GRE + TOEFL + LOR + CGPA + Research +
32                  I((GRE - mean(GRE))^2))
33 summary(poly_mod4)
34 # And last polynomial term is irrelevant
35 # i.e. we've collapsed back into original model
36
37 # This same approach is taken except starting with Y^6 instead of Y
38 poly_mod2.1 = lm(data = Admission,
39                formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA + Research +
40                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((CGPA
41                  - mean(CGPA))^2) +
42                  I((GRE - mean(GRE))^3) + I((TOEFL - mean(TOEFL))^3) + I((CGPA
43                  - mean(CGPA))^3))
44 summary(poly_mod2.1)
45 # Least relevant is (centered) TOEFL^3, remove that
46
47 poly_mod2.2 = lm(data = Admission,
48                formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA + Research +
49                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((
50                  CGPA - mean(CGPA))^2) +
51                  I((GRE - mean(GRE))^3) + I((CGPA - mean(CGPA))^3))
52 summary(poly_mod2.2)
53 # Least relevant is GRE^3, remove that

```

```

50
51 poly_mod2.3 = lm(data = Admission,
52                  formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA + Research +
53                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((
54                  CGPA - mean(CGPA))^2) +
55                  I((CGPA - mean(CGPA))^3))
56 summary(poly_mod2.3)
57 # Least relevant is GRE^2, remove that
58 poly_mod2.4 = lm(data = Admission,
59                  formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA + Research +
60                  I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) +
61                  I((CGPA - mean(CGPA))^3))
62 summary(poly_mod2.4)
63 # All terms relevant; stop here
64
65 cubic_model = lm(data = Admission,
66                  formula = I(Y^6) ~ GRE + TOEFL + LOR + CGPA + Research +
67                  I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((
68                  CGPA - mean(CGPA))^3))
69
70 plot(cubic_model)
71
72 ncvTest(cubic_model)
73 shapiro.test(cubic_model$residuals)
74
75 boxcox(cubic_model, lambda = seq(-6,6,0.1)) |> plot()
76 bc2 = boxcox(cubic_model, lambda = seq(-3,3,0.01))
77 bc2$lambda[which.max(bc2$objective)] # 0.89
78 # .89 * 6 = 5.34
79 # looks like we 'overshot' with Y^6 instead of Y^5
80 # use 5th power transform instead,
81
82 # Redo hierarchical approach, starting at Y^5 instead of Y^6
83
84 poly_mod3.1 = lm(data = Admission,
85                  formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
86                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((
87                  CGPA - mean(CGPA))^2) +
88                  I((GRE - mean(GRE))^3) + I((TOEFL - mean(TOEFL))^3) + I((
89                  CGPA - mean(CGPA))^3))
90 summary(poly_mod3.1)
91 # Drop TOEFL^3
92 poly_mod3.2 = lm(data = Admission,
93                  formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
94                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((
95                  CGPA - mean(CGPA))^2) +
96                  I((GRE - mean(GRE))^3) + I((CGPA - mean(CGPA))^3))
97 summary(poly_mod3.2)
98 # drop GRE^3
99 poly_mod3.3 = lm(data = Admission,
100                  formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
101                  I((GRE - mean(GRE))^2) + I((TOEFL - mean(TOEFL))^2) + I((
102                  CGPA - mean(CGPA))^2) +
103                  I((CGPA - mean(CGPA))^3))

```

```

98 summary(poly_mod3.3)
99 # Drop GRE^2
100 poly_mod3.4 = lm(data = Admission,
101                  formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
102                        I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) +
103                        I((CGPA - mean(CGPA))^3))
104 summary(poly_mod3.4)
105
106 shapiro.test(poly_mod3.4$residuals)
107 ncvTest(poly_mod3.4)
108 # Even still, nonnormal, heteroskedastic errors
109
110 # Further investigation done into outliers to see if omitting strong outliers
111 # can significantly improve model fit
112 # As plot show, points which weren't notable outliers in the original model
113 # may become
114 # strong outliers in new model

```

6.3.2 Identifying Outliers

```

1 p = 8
2 n = 400
3
4 # Outlying X obser.
5
6 H = hatvalues(poly_mod3.4)
7 (H > (2 * p / n)) |> which() |> unique() |> sort()
8 H |> sort(decreasing = TRUE) |> head(20)
9 #25 29 30 35 39 48 51 53 57 59 72 79 80 98 118 119 131 144 149
10    169 177 203 204 214 252 258
11 #27 285 298 345 346 348 349 369 385 386
12 # ^ strongly influential X
13 # 59 is by far strongest point here, 4x bigger h value than next highest
14
15 dff = dffits(poly_mod3.4)
16 (dff > (2*sqrt(p / n))) |> which() |> unique() |> sort()
17 # 83 287 359 360
18
19 dff[c(83, 287, 359, 360)]
20 # 360 is very influential to fitted values
21
22 cd = cooks.distance(poly_mod3.4)
23 (qf(cd, p, n-p) > 0.5) |> which() # none :)
24
25 dfb = dfbeta(poly_mod3.4)
26 (dfb > (2/sqrt(n))) |> which() # none here :)
27
28 # Per the plots, it looks like 66,67,69 are things throwing off residual vs.
29 # fitted and normal
30 # qq plot, look at Y outliers instead

```



```

30 library(MASS) # for studentized residuals
31
32 sr = studres(poly_mod3.4)
33 rejection = qt(1 - 0.05/(2*n), n - p - 1)
34 (abs(sr) > rejection) |> which() |> unique()
35 # 69 is outlier wrt. y
36
37 (abs(sr) > 3) |> which() |> unique() # using slightly lower rejection criteria
38 # 66 67 69 116 360
39 # try omitting these 5 points
40
41 cubic_mod_omitOutliers = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA +
    Research +
42                                I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))
                                ^2) + I((CGPA - mean(CGPA))^3),
43                                data = Admission[-c( 66, 67 ,69, 116, 360), ])
44
45 shapiro.test(cubic_mod_omitOutliers$residuals) # 0.09864
46 ncvTest(cubic_mod_omitOutliers) # 0.00068006
47 # Heteroskedasticity still present, but residuals are now reasonably normal
48 # weighting can now be applied to model

```

6.4 Weighted Model

```

1 res = cubic_mod_omitOutliers$residuals
2 fitted = cubic_mod_omitOutliers$fitted.values
3 mod = lm(formula = abs(res) ~ fitted)
4
5 var = mod$fitted.values^2
6 w.i = 1/var
7
8 cubic.wls = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
9                I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((
10                  CGPA - mean(CGPA))^3),
11                data = Admission[-c( 66, 67 ,69, 116, 360), ],
12                weights = w.i)
13 summary(cubic.wls)
14 plot(cubic.wls)
15 shapiro.test(cubic.wls$residuals) # 0.2
16 ncvTest(cubic.wls) # 0.41676
17 # So, this weighted cubic model does meet assumptions of regression :)

```

6.5 Bootstrapping

6.5.1 Confidence Intervals

```

1   # since the family wise confidence level will be 90%, each CI will use
2   # alpha = 0.1/9 = 0.0111... , i.e. each CI is at 98.889% level
3
4   set.seed(123)
5
6   sample_from = setdiff(1:400, c( 66, 67 ,69, 116, 360))
7   # only want to sample from the data used to construct model
8
9   trials = 1000
10  coeff_values = data.frame(matrix(nrow = 1000, ncol = 9))
11  colnames(coeff_values) = c("b0", "b1", "b2", "b3", "b4", "b5", "b22", "b44", "b444")
12
13  for(i in 1:trials){
14    # Generate sample
15    indices = sample(sample_from, size = 395, replace = T)
16    bt_samp = Admission[indices,]
17
18    # First, fit unweighted model
19    mod1 = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
20              I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((CGPA
21                - mean(CGPA))^3),
22              data = bt_samp)
23
24    # Use that to generate weights
25    res = mod1$residuals
26    fitted = mod1$fitted.values
27    mod = lm(formula = abs(res) ~ fitted)
28
29    var = mod$fitted.values^2
30    sim_weights = 1/var
31
32    # Get coefficient estimates of weighted model
33    sim_mod = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
34                 I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((
35                   CGPA - mean(CGPA))^3),
36                 data = bt_samp, weights = sim_weights)
37
38    # Record these coefficients
39    coeff_values[i,] = coefficients(sim_mod)
40    # and repeat :)
41  }
42
43  write.csv(coeff_values, "./bootstrap_coefficients", row.names = F)
44
45  library(ggplot2)
46  # Generate histograms for each coefficient
47
48  plot_b0 = ggplot(data = coeff_values, aes(x = b0)) +
49    theme_classic() +
50    geom_histogram(bins = 50) +
51    ylab("") +
52    labs(title = "Simulated b0 values") +
53    xlab("") +

```

```

52 scale_y_continuous(expand = expansion(mult = 0)) +
53 theme(plot.title = element_text(size = 20, hjust = 0.5),
54       axis.text.x = element_text(size = 15),
55       axis.text.y = element_text(size = 15),
56       axis.title.x = element_text(size = 15),
57       axis.title.y = element_text(size = 15))
58
59 plot_b1 = ggplot(data = coeff_values, aes(x = b1)) +
60 theme_classic() +
61 geom_histogram(bins = 50) +
62 ylab("") +
63 labs(title = "Simulated b1 values") +
64 xlab("") +
65 scale_y_continuous(expand = expansion(mult = 0)) +
66 theme(plot.title = element_text(size = 20, hjust = 0.5),
67       axis.text.x = element_text(size = 15),
68       axis.text.y = element_text(size = 15),
69       axis.title.x = element_text(size = 15),
70       axis.title.y = element_text(size = 15))
71
72
73 plot_b2 = ggplot(data = coeff_values, aes(x = b2)) +
74 theme_classic() +
75 geom_histogram(bins = 50) +
76 ylab("") +
77 labs(title = "Simulated b2 values") +
78 xlab("") +
79 scale_y_continuous(expand = expansion(mult = 0)) +
80 theme(plot.title = element_text(size = 20, hjust = 0.5),
81       axis.text.x = element_text(size = 15),
82       axis.text.y = element_text(size = 15),
83       axis.title.x = element_text(size = 15),
84       axis.title.y = element_text(size = 15))
85
86
87 plot_b3 = ggplot(data = coeff_values, aes(x = b3)) +
88 theme_classic() +
89 geom_histogram(bins = 50) +
90 ylab("") +
91 labs(title = "Simulated b3 values") +
92 xlab("") +
93 scale_y_continuous(expand = expansion(mult = 0)) +
94 theme(plot.title = element_text(size = 20, hjust = 0.5),
95       axis.text.x = element_text(size = 15),
96       axis.text.y = element_text(size = 15),
97       axis.title.x = element_text(size = 15),
98       axis.title.y = element_text(size = 15))
99
100
101 plot_b4 = ggplot(data = coeff_values, aes(x = b4)) +
102 theme_classic() +
103 geom_histogram(bins = 50) +
104 ylab("") +
105 labs(title = "Simulated b4 values") +

```

```

106 xlab("") +
107 scale_y_continuous(expand = expansion(mult = 0)) +
108 theme(plot.title = element_text(size = 20, hjust = 0.5),
109       axis.text.x = element_text(size = 15),
110       axis.text.y = element_text(size = 15),
111       axis.title.x = element_text(size = 15),
112       axis.title.y = element_text(size = 15))
113
114
115 plot_b5 = ggplot(data = coeff_values, aes(x = b5)) +
116   theme_classic() +
117   geom_histogram(bins = 50) +
118   ylab("") +
119   labs(title = "Simulated b5 values") +
120   xlab("") +
121   scale_y_continuous(expand = expansion(mult = 0)) +
122   theme(plot.title = element_text(size = 20, hjust = 0.5),
123         axis.text.x = element_text(size = 15),
124         axis.text.y = element_text(size = 15),
125         axis.title.x = element_text(size = 15),
126         axis.title.y = element_text(size = 15))
127
128
129 plot_b22 = ggplot(data = coeff_values, aes(x = b22)) +
130   theme_classic() +
131   geom_histogram(bins = 50) +
132   ylab("") +
133   labs(title = "Simulated b22 values") +
134   xlab("") +
135   scale_y_continuous(expand = expansion(mult = 0)) +
136   theme(plot.title = element_text(size = 20, hjust = 0.5),
137         axis.text.x = element_text(size = 15),
138         axis.text.y = element_text(size = 15),
139         axis.title.x = element_text(size = 15),
140         axis.title.y = element_text(size = 15))
141
142 plot_b44 = ggplot(data = coeff_values, aes(x = b44)) +
143   theme_classic() +
144   geom_histogram(bins = 50) +
145   ylab("") +
146   labs(title = "Simulated b44 values") +
147   xlab("") +
148   scale_y_continuous(expand = expansion(mult = 0)) +
149   theme(plot.title = element_text(size = 20, hjust = 0.5),
150         axis.text.x = element_text(size = 15),
151         axis.text.y = element_text(size = 15),
152         axis.title.x = element_text(size = 15),
153         axis.title.y = element_text(size = 15))
154
155 plot_b444 = ggplot(data = coeff_values, aes(x = b444)) +
156   theme_classic() +
157   geom_histogram(bins = 50) +
158   ylab("") +
159   labs(title = "Simulated b444 values") +

```

```

160 xlab("") +
161 scale_y_continuous(expand = expansion(mult = 0)) +
162 theme(plot.title = element_text(size = 20, hjust = 0.5),
163        axis.text.x = element_text(size = 15),
164        axis.text.y = element_text(size = 15),
165        axis.title.x = element_text(size = 15),
166        axis.title.y = element_text(size = 15))
167
168
169 library(ggpubr) # for ggarrange
170
171 ggarrange(plot_b0, plot_b1, plot_b2, plot_b3, plot_b4, plot_b5, plot_b22, plot_b44,
172           plot_b444)
173
174 # Get quantile values
175 family_alpha = 0.1
176 alpha = family_alpha/9
177
178 quantile(coeff_values["b0"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
179 quantile(coeff_values["b1"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
180 quantile(coeff_values["b2"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
181 quantile(coeff_values["b3"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
182 quantile(coeff_values["b4"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
183 quantile(coeff_values["b5"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
184 quantile(coeff_values["b22"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round(4)
185 quantile(coeff_values["b44"][,1], probs = c(alpha/2, 1 - alpha/2)) |> round
  (4)

```

6.5.2 Prediction Intervals

```

1 h = hatvalues(cubic.wls)
2 min(h) # 0.006827814
3 max(h) # 0.4296308
4 # ^ range of acceptable values for making predictions and ensure no
  extrapolation
5
6 # See how presence of research affects admission chance at different levels of
  other variables
7 # Average Grades/test scores
8 to_predict1.1 = data.frame("GRE" = 320, "TOEFL" = 110, "LOR" = 4.0, "CGPA" =
  8.6, "Research" = 0)
9 to_predict1.2 = data.frame("GRE" = 320, "TOEFL" = 110, "LOR" = 4.0, "CGPA" =
  8.6, "Research" = 1)
10 # Poor Grades/test scores
11 to_predict2.1 = data.frame("GRE" = 310, "TOEFL" = 105, "LOR" = 3.0, "CGPA" =
  7.5, "Research" = 0)
12 to_predict2.2 = data.frame("GRE" = 310, "TOEFL" = 105, "LOR" = 3.0, "CGPA" =
  7.5, "Research" = 1)
13 # Excellent Grades/test scores

```

```

14 to_predict3.1 = data.frame("GRE" = 330, "TOEFL" = 113, "LOR" = 5.0, "CGPA" =
    9.8, "Research" = 0)
15 to_predict3.2 = data.frame("GRE" = 330, "TOEFL" = 113, "LOR" = 5.0, "CGPA" =
    9.8, "Research" = 1)
16
17 # Check to make sure any predictions aren't extrapolation
18
19 X = model.matrix(cubic.wls) |> unname()
20
21 obs1 = matrix(data = c(1,320, 110, 4.0,8.61, 0, (110 - 107.3823)^2,
    (8.6-8.594759)^2, (8.6-8.594759)^3),
22               ncol = 1)
23 obs2 = matrix(data = c(1,320, 110, 4.0,8.61, 1, (110 - 107.3823)^2,
    (8.6-8.594759)^2, (8.6-8.594759)^3),
24               ncol = 1)
25 obs3 = matrix(data = c(1,310, 105, 3.0,7.5, 0, (105 - 107.3823)^2,
    (7.5-8.594759)^2, (7.5-8.594759)^3),
26               ncol = 1)
27 obs4 = matrix(data = c(1,310, 105, 3.0,7.5, 1, (105 - 107.3823)^2,
    (7.5-8.594759)^2, (7.5-8.594759)^3),
28               ncol = 1)
29 obs5 = matrix(data = c(1,330, 113, 5.0,9.8, 0, (113 - 107.3823)^2,
    (9.8-8.594759)^2, (9.8-8.594759)^3),
30               ncol = 1)
31 obs6 = matrix(data = c(1,330, 113, 5.0,9.8, 1, (113 - 107.3823)^2,
    (9.8-8.594759)^2, (9.8-8.594759)^3),
32               ncol = 1)
33
34 t(obs1) %*% (t(X) %*% X)^-1 %*% obs1 # 0.07276908
35 t(obs2) %*% (t(X) %*% X)^-1 %*% obs2 # 0.1252422
36 t(obs3) %*% (t(X) %*% X)^-1 %*% obs3 # -1.318934
37 t(obs4) %*% (t(X) %*% X)^-1 %*% obs4 # -1.300921
38 t(obs5) %*% (t(X) %*% X)^-1 %*% obs5 # 3.418855
39 t(obs6) %*% (t(X) %*% X)^-1 %*% obs6 # 3.596004
40
41 # only first 2 points not extrapolation, just look at those
42
43 to_predict1.1 = data.frame("GRE" = 320, "TOEFL" = 110, "LOR" = 4.0, "CGPA" =
    8.6, "Research" = 0)
44 to_predict1.2 = data.frame("GRE" = 320, "TOEFL" = 110, "LOR" = 4.0, "CGPA" =
    8.6, "Research" = 1)
45
46
47 set.seed(777) # for reproducibility
48
49 trials = 5000
50
51 predicted_values = data.frame(matrix(nrow = 5000, ncol = 2))
52 colnames(predicted_values) = c("Research", "No Research")
53
54 for(i in 1:trials){
55   # Generate sample
56   indices = sample(sample_from, size = 395, replace = T)
57   bt_samp = Admission[indices,]

```

```

58
59 # First, fit unweighted model
60 mod1 = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
61           I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((CGPA
62             - mean(CGPA))^3),
63           data = bt_samp)
64
65 # Use that to generate weights
66 res = mod1$residuals
67 fitted = mod1$fitted.values
68 mod = lm(formula = abs(res) ~ fitted)
69
70 var = mod$fitted.values^2
71 sim_weights = 1/var
72
73 sim_mod = lm(formula = I(Y^5) ~ GRE + TOEFL + LOR + CGPA + Research +
74             I((TOEFL - mean(TOEFL))^2) + I((CGPA - mean(CGPA))^2) + I((
75               CGPA - mean(CGPA))^3),
76             data = bt_samp, weights = sim_weights)
77
78 # make predictions with simulated model & record them
79 predicted_values[i,1] = predict(sim_mod, to_predict1.2) # with research
80 predicted_values[i,2] = predict(sim_mod, to_predict1.1) # w/o research
81 # repeat :)
82 }
83
84 write.csv(predicted_values, "./bt-prediction", row.names = F)
85
86 library(ggplot2)
87
88 # Generate histograms for each prediction
89 plot_no_research = ggplot(data = predicted_values, aes(x = 'No Research'^(1/5)
90 )) +
91   theme_classic() +
92   geom_histogram(bins = 100) +
93   ylab("Frequency") +
94   labs(title = "Predictions for no undergrad research") +
95   xlab("Reverse-transformed probability") +
96   scale_y_continuous(expand = expansion(mult = 0)) +
97   theme(plot.title = element_text(size = 20, hjust = 0.5),
98         axis.text.x = element_text(size = 15),
99         axis.text.y = element_text(size = 15),
100        axis.title.x = element_text(size = 15),
101        axis.title.y = element_text(size = 15))
102
103 plot_research = ggplot(data = predicted_values, aes(x = Research^(1/5))) +
104   theme_classic() +
105   geom_histogram(bins = 100) +
106   ylab("Frequency") +
107   labs(title = "Predictions for completed undergrad research") +
108   xlab("Reverse-transformed probability") +
109   scale_y_continuous(expand = expansion(mult = 0)) +
110   theme(plot.title = element_text(size = 20, hjust = 0.5),

```

```

109     axis.text.x = element_text(size = 15),
110     axis.text.y = element_text(size = 15),
111     axis.title.x = element_text(size = 15),
112     axis.title.y = element_text(size = 15))
113
114
115
116 # Want the 2 PI's at 95% level, family wide
117 # PIs are calculated on transformed data and endpoints are then translated
118   back into
119 # original terms
120 family_alpha = 0.05
121 alpha = family_alpha/2 #0.025
122
123 quantile(predicted_values$Research, probs = c(alpha/2, 1 - alpha / 2))
124 # 0.2471215 0.2819776
125 quantile(predicted_values$'No Research', probs = c(alpha/2, 1 - alpha / 2))
126 # 0.2128977 0.2486968
127
128 # in terms of original data
129 (c(0.2471215, 0.2819776)^(1/5)) |> round(4) # 0.7561 0.7763
130 (c(0.2128977, 0.2486968)^(1/5)) |> round(4) # 0.7339 0.7571
131
132 ggarrange(plot_research, plot_no_research, nrow = 2)

```